

CS 6290 Project 3: Cache Coherence

Sam Britt

April 27, 2012

Results

Figure shows the runtime (in clock cycles) of the trace simulations for each cache coherency protocol. Because of the order of magnitude difference in runtime, the longer traces used for validation are plotted on the left, and the eight shorter traces are plotted separately on the right. (All results will be presented this way).

Figures and show the number of cache misses and cache-to-cache transfers, respectively, that occurred during the simulation of each trace for each coherency protocol. Figure shows, for each trace, the number of silent upgrades from the exclusive to the modified state. Data only those protocols that have an exclusive state, namely MESI, MOESI, and MOESIF, are shown.




Figure 1: Runtime (in clock cycles) of the various trace simulations using each cache coherency protocol. The longer validation runs are on the left plot, and the short traces are on the right.

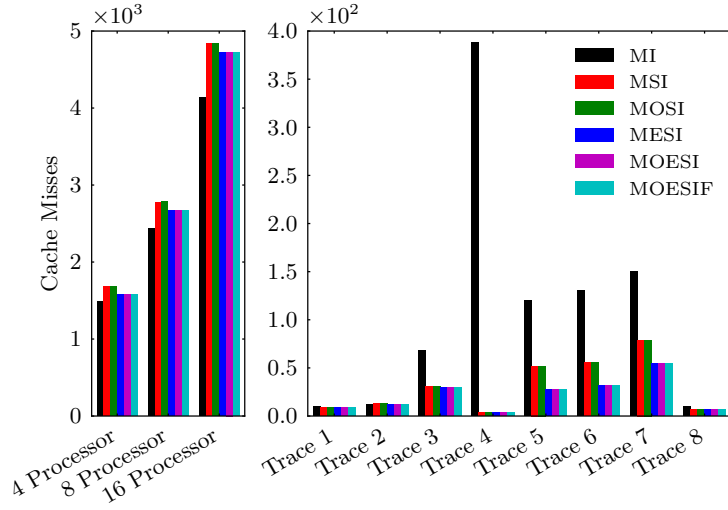


Figure 2: Number of cache misses during simulation of the various traces using each cache coherence protocol.

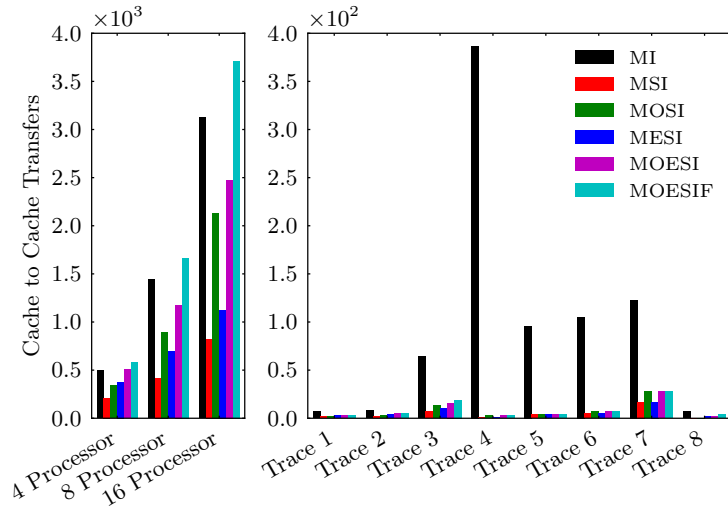


Figure 3: Number of cache-to-cache transfers during simulation of the various traces using each cache coherence protocol.




Figure 4: Number of silent upgrades to the modified state from the exclusive state during simulation of the various traces using each cache coherency protocol that contains an exclusive state.

Discussion

Interestingly, the simple MI protocol outperformed most or all other protocols in almost every trace, as can be seen from Figure . Compare with the runtime results for the MSI protocol: with the exception of trace 4 (which will be discussed later), adding the shared state to the protocol increased runtime by a factor of as much as 4. This is perhaps not as expected, since adding a shared state allows multiple caches to keep read-only copies of data, which should reduce the number of cache misses; indeed this is the case, as shown by Figure . In every simulation of the short traces, the number of cache misses is reduced by sharing, and the effect is particularly evident for the read-heavy traces 3–7.

However, adding the shared state removes one critical characteristic of the simpler MI protocol: the high degree of cache-to-cache transfers. Even if a processor needs only read-only accesses of a cache line, in the MI protocol it still must be put in the modified state, and thus it will supply the data over the bus on the next request, avoiding the costly trip to memory. The effect is quite clearly shown by Figure . The MI protocol may not allow true sharing, but sharing may not be necessary for performance if the bus has sufficiently low latency.

One interesting trace was trace 4. It seems to buck the trend: the MSI protocol outperformed the MI protocol, and the number of cache misses and cache-to-cache transfers using the MI protocol dominated all the other traces and protocols. Upon inspection of the trace itself the reason became apparent.

Trace 4 simulated a machine with four processors. After an initial write to a particular memory address by the first processor, all accesses from all processors are reads to that address. Clearly, such an access pattern benefits from the sharing provided by the MSI protocol. The penalty of going to memory is avoided because, after the initial access from each processor, the data in the cache is clean and can be provided to the processor directly—no need for bus or memory interaction. So applications with such a read-heavy access pattern would benefit from such a protocol.

Though most traces performed well in the MOESI and MOESIF protocols, traces 5, 6, and 7 seemed to particularly benefit from the addition of the exclusive state. A quick look at Figure shows the reason: the access patterns of those traces allowed for a larger number of silent upgrades from the exclusive state to the modified state. This reduced the number of cache misses (see Figure) since, without the exclusive state, such a cache line would be in the shared state and would have to go to the bus to get modified access.

Conclusion

Based on the results of the simulations, there is little reason to use anything other than the MI cache coherency protocol. In many cases, adding complexity to the protocol hurt performance, and in the other cases, the improvement in performance was not so much to overwhelm the increase in cost and complexity of the implementing the other protocols. The success of the MI protocol, though, is largely due to the high discrepancy between the bus and memory controller latencies. If that gap were closed (by perhaps a faster memory controller, a slower bus), or if there were many more processors contending for the bus, the simpler protocol may not be sufficient. Also, if the memory access pattern is known *a priori* to be, say almost exclusively read-only, a protocol may be chosen that optimizes for that use case.